



PARALLEL **P**ROGRAMMING...

Copyright 2023 Patrick Lemoine. All rights reserved.

Parallel Programming: Overview

SESSION 2/3



Programming Interface for parallel computing

MPI (Message Passing Interface)

OpenMP (Open Multi-Processing)

CUDA and ROCm

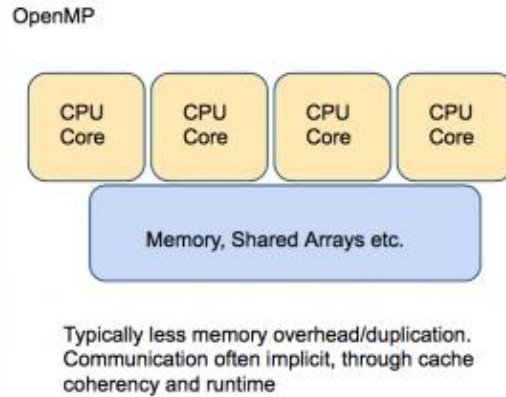

Programming **i**nterface for **p**arallel **c**omputing

병렬 컴퓨팅을 위한 프로그래밍
인터페이스

Programming interface...



Remember



MPI (Message Passing Interface) is a multi-process model whose mode of communication between the processes is **explicit**.

==> communication management is the responsibility of the user.

OpenMP (Open Multi-Processing) is a multitasking model whose mode of communication between tasks is **implicit**

==> communications is the responsibility of the compiler.



MPI (**M**essage **P**assing **I**nterface)



MPI (Message Passing Interface)

MPI is a library of subroutines (in Fortran,C)

MPI allows the coordination of a program running as multiple processes in a distributed-memory environment.

MPI is flexible enough to also be used in a shared-memory environment.

MPI programs can be used and compiled on a wide variety of single platforms or (homogeneous or heterogeneous) clusters of computers over a network.

MPI library is standardized, should work (without further changes!) on any machine on which the MPI library is installed.



MPI: Basic Environment



```
MPI_Init(&argc, &argv)
```

Initializes MPI environment

Must be called in every MPI program

Must be first MPI call

Can be used to pass command line arguments to all

```
MPI_Finalize()
```

Terminates MPI environment

Last MPI function call

MPI: Basic Environment



```
MPI_Comm_rank(comm, &rank)
```

Returns the rank of the calling MPI process

Within the communicator, comm

MPI_COMM_WORLD is set during Init(...)

Other communicators can be created if needed

```
MPI_Comm_size(comm, &size)
```

Returns the total number of processes

Within the communicator, comm

```
int my_rank, size;  
MPI_Init(&argc, &argv);  
MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);  
MPI_Comm_size(MPI_COMM_WORLD, &size);
```


MPI: Point-to-Point Communication



```
MPI_Send(&buf, count, datatype, dest, tag, comm)
```

Send a message

Returns only after buffer is free for reuse (Blocking)

```
MPI_Recv(&buf, count, datatype, source, tag, comm, &status)
```

Received a message


Returns only when the data is available

Blocking

```
MPI_SendRecv(...)
```

Two way communication

Blocking



MPI: Point-to-Point Communication



Blocking

- Only returns after completed
 - Received: data has arrived and ready to use
 - Send: safe to reuse sent buffer
- Be aware of deadlocks
- Tip: use when possible



Non-Blocking

- returns immediately
 - Unsafe to modify buffers until operation is known to be complete
- Allows computation and communication to overlap
- Tip: Use only when needed

MPI: Deadlock

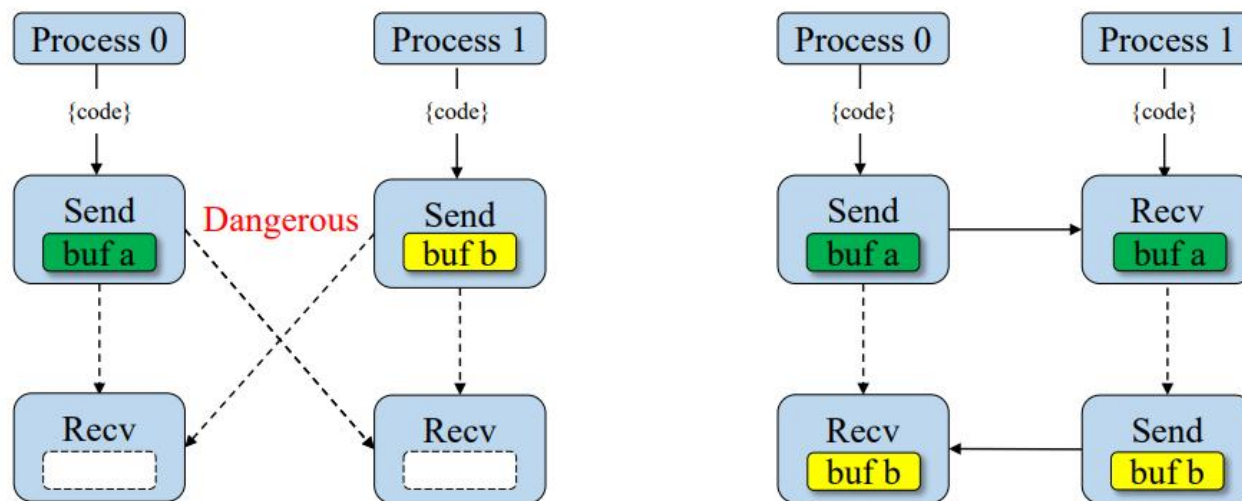


Blocking calls can result in deadlock

One process is waiting for message that will never arrive

Only option is to abort the interrupt/kill the code (CTRL-c)

Might not always deadlock - depends on size of system buffer




MPI: Coolective Communication (BCast)



```
MPI_Bcast(&buffer, count, datatype, root, comm)
```

Broadcast a message from the root process to all other processes.
Useful when reading in input parameters from file.




MPI: Collective Communication (Scatter)



```
MPI_Scatter(&sendbuf, sendcnt, sendtype, &recvbuf,  
           recvcnt, recvtype, root, comm)
```

Sends individual messages from the root process to all other processes.




MPI: Collective Communication (Gather)



```
MPI_Gather(&sendbuf, sendcnt, sendtype, &recvbuf,  
          recvnt, recvtype, root, comm)
```

Opposite of Scatter.




MPI: Collective Communication (Reduce)



```
MPI_Reduce(&sendbuf, &recvbuf, count, datatype,  
           mpi_operation, root, comm)
```

Applies reduction operation on data from all processes.
Puts results on root process.

| Operator |
|----------|
| MPI_SUM |
| MPI_MAX |
| MPI_MIN |
| MPI_PROD |




MPI: Collective Communication (Allreduce)



```
MPI_Allreduce(&sendbuf, &recvbuf, count,  
             datatype, mpi_operation, comm)
```

Applies reduction operation on data from all processes.
Store results on all processes.

| Operator |
|----------|
| MPI_SUM |
| MPI_MAX |
| MPI_MIN |
| MPI_PROD |



MPI: Collective Communication (Barrier)



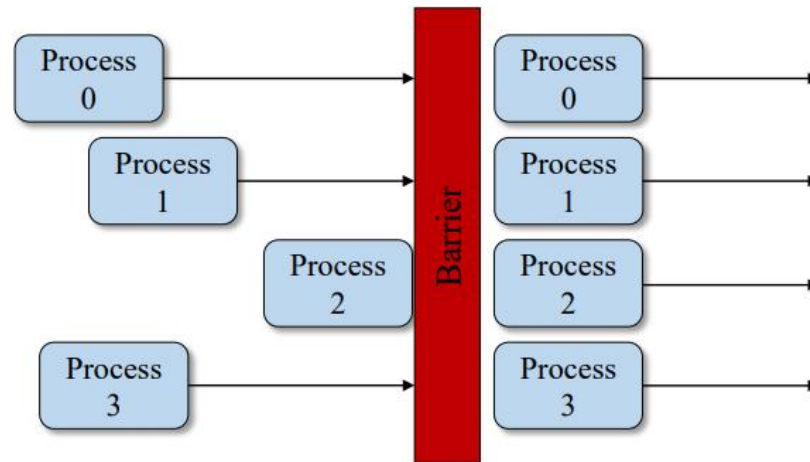
```
MPI_Barrier(comm)
```

Process synchronization (blocking).

All processes forced to wait for each other.

Use only where necessary.

Will reduce parallelism.



MPI: keywords

1 environment

- MPI Init: Initialization of the MPI environment
- MPI Comm rank: Rank of the process
- MPI Comm size: Number of processes
- MPI Finalize: Deactivation of the MPI environment
- MPI Abort: Stopping of an MPI program
- MPI Wtime: Time taking

2 Point-to-point communications

- MPI Send: Send message
- MPI Isend: Non-blocking message sending
- MPI Recv: Message received
- MPI Irecv: Non-blocking message reception
- MPI Sendrecv and MPI Sendrecv replace: Sending and receiving messages
- MPI Wait: Waiting for the end of a non-blocking communication
- MPI Wait all: Wait for the end of all non-blocking communications

3 Collective communications

- MPI Bcast: General broadcast
- MPI Scatter: Selective spread
- MPI Gather and MPI Allgather: Collecting
- MPI Alltoall: Collection and distribution
- MPI Reduce and MPI Allreduce: Reduction
- MPI Barrier: Global synchronization



4 Derived Types

- MPI Contiguous type: Contiguous types
- MPI Type vector and MPI Type create hvector: Types with a constant
- MPI Type indexed: Variable pitch types
- MPI Type create subarray: Sub-array types
- MPI Type create struct: H and erogenous types
- MPI Type commit: Type commit
- MPI Type get extent: Recover the extent
- MPI Type create resized: Change of scope
- MPI Type size: Size of a type
- MPI Type free: Release of a type

MPI: Keywords

5 Communicator

- MPI Comm split: Partitioning of a communicator
- MPI Dims create: Distribution of processes
- MPI Cart create: Creation of a Cartesian topology
- MPI Cart rank: Rank of a process in the Cartesian topology
- MPI Cart coordinates: Coordinates of a process in the Cartesian topology
- MPI Cart shift: Rank of the neighbors in the Cartesian topology
- MPI Comm free: Release of a communicator

6 MPI-IO

- MPI File open: Opening a file
- MPI File set view: Changing the view
- MPI File close: Closing a file

6.1 Explicit addresses

- MPI File read at: Reading
- MPI File read at all: Collective reading
- MPI File write at: Writing

6.2 Individual pointers

- MPI File read: Reading
- MPI File read all: collective reading
- MPI File write: Writing
- MPI File write all: collective writing
- MPI File seek: Pointer positioning

6.3 Shared pointers

- MPI File read shared: Read
- MPI File read ordered: Collective reading
- MPI File seek shared: Pointer positioning

7.0 Symbolic constants

- MPI COMM WORLD, MPI SUCCESS
- MPI STATUS IGNORE, MPI PROC NULL
- MPI INTEGER, MPI REAL, MPI DOUBLE PRECISION
- MPI ORDER FORTRAN, MPI ORDER C
- MPI MODE CREATE, MPI MODE ONLY, MPI MODE EXIST, MPI MODE EXISTING



MPI: Program Basics



Include MPI Header File

Start of Program
(Non-interacting Code)

Initialize MPI

Run Parallel Code &
Pass Messages

End MPI Environment

(Non-interacting Code)

End of Program

```
#include <mpi.h>
```

```
int main (int argc, char *argv[])  
{
```

```
    MPI_Init(&argc, &argv);
```

```
    .  
    .    // Run parallel code  
    .
```

```
    MPI_Finalize(); // End MPI Envir
```

```
    return 0;  
}
```

MPI: Example



```
#include <mpi.h>
#include <stdio.h>

int main (int argc, char *argv[]) {

    int rank, size;

    MPI_Init (&argc, &argv); //initialize MPI library

    MPI_Comm_size(MPI_COMM_WORLD, &size); //get number of processes
    MPI_Comm_rank(MPI_COMM_WORLD, &rank); //get my process id

    //do something
    printf ("Hello World from rank %d\n", rank);
    if (rank == 0) printf("MPI World size = %d processes\n", size);

    MPI_Finalize(); //MPI cleanup

    return 0;
}
```

- 4 processes

```
Hello World from rank 3
Hello World from rank 0
MPI World size = 4 processes
Hello World from rank 2
Hello World from rank 1
```

- Code ran on each process independently
- MPI Processes have *private* variables
- Processes can be on completely different machines

COMPILING an MPI Program



- **Compiling a program** for MPI is almost just like compiling a regular C or C++ program
 - The C compiler is **mpicc** and the C++ compiler is **mpic++**.
 - For example, to compile **MyProg.c** you would use a command like
 - **mpicc -O2 -o MyProg MyProg . c**





OpenMP (Open Multi-Processing)

OpenMP



OpenMP (Open Multi-Processing)

OpenMP is a programming interface for parallel computing on shared memory architecture.

- **It allows you to manage:**



- the creation of light processes
- the sharing of work between these lightweight processes
- synchronizations (explicit or implicit) between all light processes
- the status of the variables (private or shared).

OpenMP (Open Multi-Processing)

- **Shared memory model**

- Threads communicate by accessing shared variables

- **The sharing is defined syntactically**

- Any variable that is seen by two or more threads is shared
- Any variable that is seen by one thread only is private



- **Race conditions possible**

- Use synchronization to protect from conflicts
- Change how data is stored to minimize the synchronization

OpenMP (Open Multi-Processing)

- **Multicore CPUs are everywhere:**

- Servers with over 100 cores today and more
- Even smartphone CPUs have 8 cores

- **Multithreading, natural programming model**

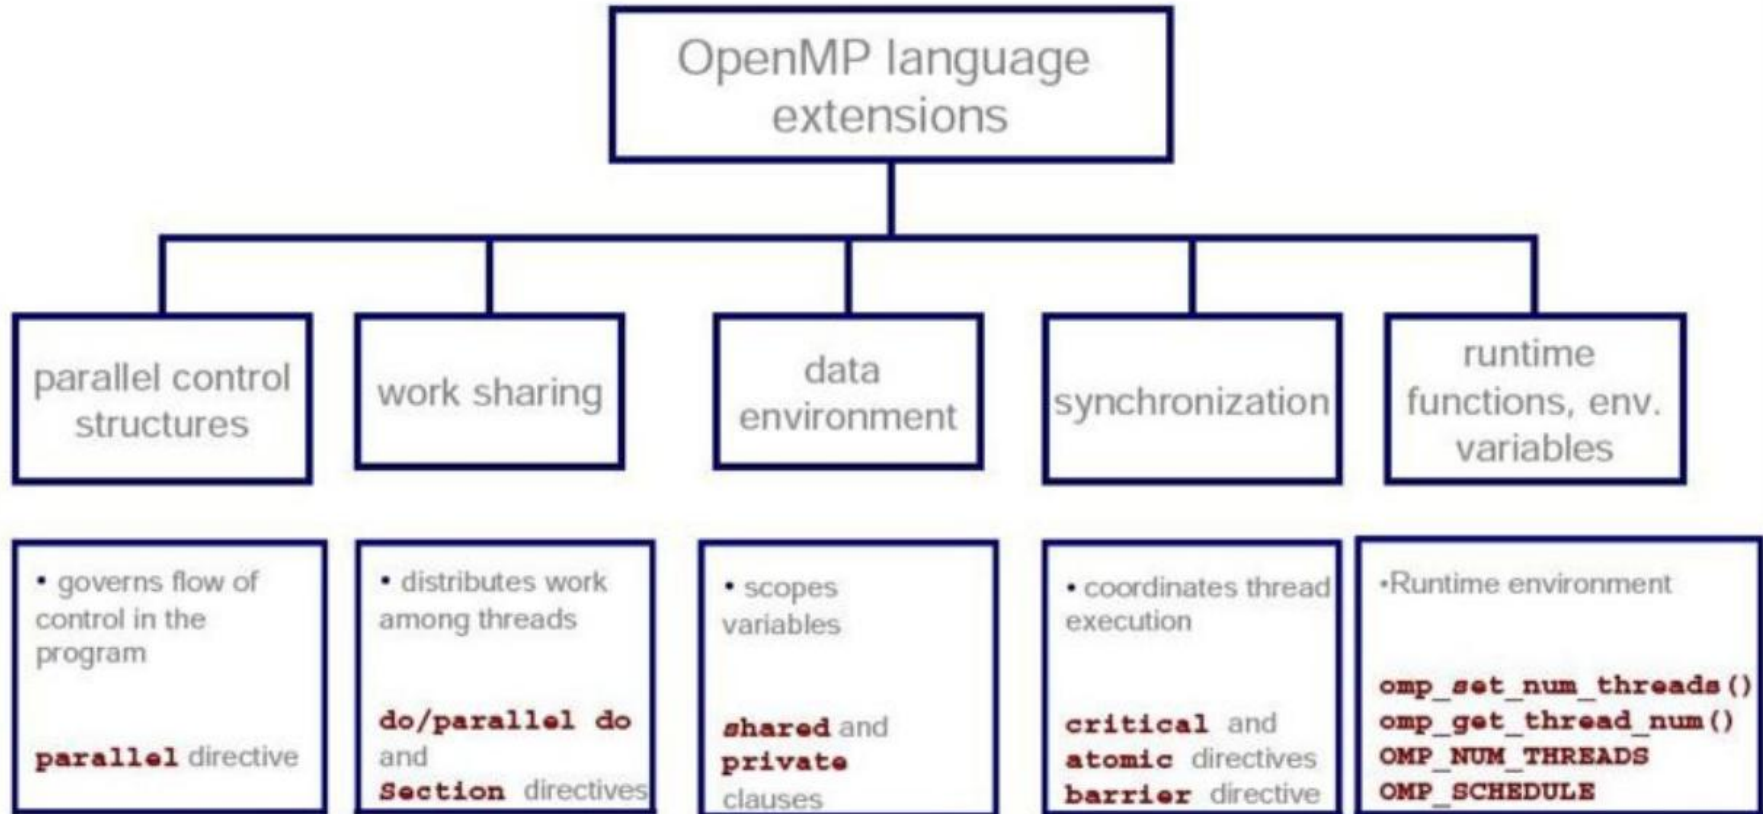
- All processors share the same memory
- Threads in a process see same address space
- Many shared-memory algorithms developed



- **Multithreading is hard**


- Lots of expertise necessary
- Deadlocks and race conditions
- **Non-deterministic** behavior makes it hard to debug

OpenMP: Architecture




OpenMP: Terminology and behavior

- **OpenMP Team** = Master + Worker
- **Parallel Region** is a block of code executed by all threads simultaneously (**has implicit barrier**)
 - The master thread always has thread id 0
 - Parallel regions can be nested
 - If clause can be used to guard the parallel region



OpenMP: Terminology and behavior

- A **Work-Sharing construct** divides the execution of the enclosed code region among the members of the team. (Loop, Section etc.)



OpenMP: Preprocessor Directives



Preprocessor directives tell the compiler what to do.

Always start with #

You have already seen one:

```
#include <stdio.h>
```

OpenMP directives tell the compiler to add machine code for parallel execution of the following block.

```
#pragma omp parallel
```

OpenMP: Some OpenMP Subroutines



```
int omp_get_max_threads()
```

Return max possible

```
int omp_get_num_threads()
```

Returns number of threads in current team \\\

```
int omp_get_thread_num()
```

Returns thread id of calling thread

Between () and omp_get_num_threads-1

OpenMP: Process vs. Thread



- MPI = Process, OpenMP = Thread
- Program start with a single process
- Process have their own (private) memory space
- A process can create one or more threads
- Threads created by a process share its memory space
 - Read and write to same memory addresses
 - Share same process ids and file descriptors
- Each thread has a unique counter and stack pointer
 - A tread can have private starage on the stack



OpenMP: Example



```
#include <omp.h> //<-- necessary header file for OpenMP API
#include <stdio.h>

int main(int argc, char *argv[]){

    printf("OpenMP running with %d threads\n", omp_get_max_threads());

#pragma omp parallel
    {
        //Code here will be executed by all threads
        printf("Hello World from thread %d\n", omp_get_thread_num());
    }

    return 0;
}
```

OpenMP: Constructs



Parallel region

Thread creates team, and becomes master (id 0)

All threads run code after

Barrier at end of parallel section

```
#pragma omp parallel [clause ...]
    if (scalar_expression)
    private (list)
    shared (list)
    default (shared | none)
    firstprivate (list)
    lastprivate (list)
    reduction (operator: list)
    num_threads (integer)
```

structured_block

(not a complete list)

OpenMP Parallel Clauses



```
#pragma omp parallel if (scalar_expression)
```

Only execute in parallel.
Otherwise serial.

```
#pragma omp parallel private (list)
```

Data local to thread.
Value are **not guaranteed to be defined on exit** (even if defined before)
No storage associated with original object
Use firstprivate and/or lastprivate clause to override

```
#pragma omp parallel firstprivate (list)
```

Variables in list are private.
Initialized with the value the variable had before entering the construct.

```
#pragma omp parallel for lastprivate (list)
```

Only in for loops
Variables in list are private.
The thread that executes the sequentially last iteration updates the value of the variables in the list.

OpenMP Parallel Clause 3



```
#pragma omp shared (list)
```

Data is accessible by all threads in team.
All threads access samme address space.

Improperly scoped variables are big source of OMP bugs

- Shared when should be private
- Race condition

```
#pragma omp default (shared | none)
```

Tip: Safest is to use default (none) and declare by hand.


OpenMP: Caution Race Condition



When multiple threads simultaneously read/write
Multiple OMP solutions

- Reduction
- Atomic
- Critical

```
#pragma omp parallel for private(i) shared(sum)
for (i=0; i<N; i++) {
    sum += i;
}
```



Should be 3!


OpenMP: Critical Section



One solution: use critical
Only one thread at a time can execute
a critical section

```
#pragma omp critical
{
    sum += i;
}
```

Downside ?
SLOOOOWWW
Overhead and serialization




OpenMP: Atomic



Atoms like "mini" critical
Only one line
Certain limitations

```
#pragma omp atomic  
sum += i;
```

Hardware controlled
Less overhead than critical



OpenMP Reduction



```
#pragma omp reduction (operator:variable)
```

Avoids race condition

Reduce variable must be shared

Makes variable private, then performs operator at end of loop

operator cannot be overloaded (c++)

One of: +, *, -, / (and &, ^, |, &&, ||)

OpenMP 3.1: added min and max for c/c++

```
#include <omp.h>
#include <stdio.h>

int main() {

    int i;
    const int N = 1000;
    int sum = 0;

    #pragma omp parallel for private(i) reduction(+: sum)
    for (i=0; i<N; i++) {
        sum += i;
    }

    printf("reduction sum=%d (expected %d)\n", sum, ((N-1)*N)/2);
}
```


OpenMP: Scheduling omp for



How does a loop get split up ?

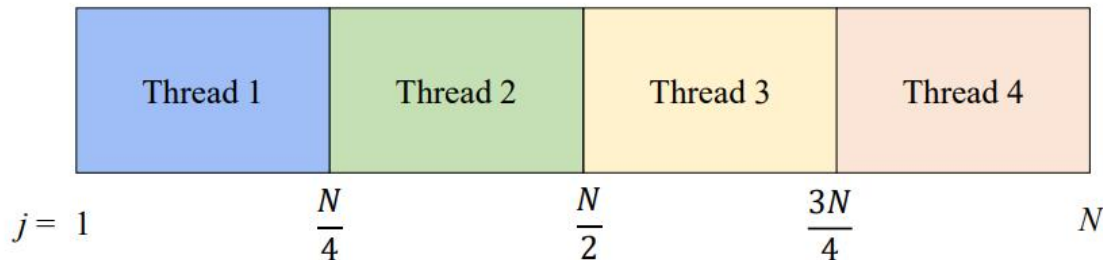
In MPI, we have to do it manually

If you do not tell what to do, the compiler decides

Usually compiler chooses "static" - chunks of N/p

```
#pragma omp parallel for default(shared) private(j)
  for (j=0; j<N; j++) {
    ... // some work here
  }
```

Unspecified schedule




OpenMP: Static Scheduling



You can tell the compiler what size chunks to take.

```
#pragma omp parallel for default(shared) private(j) schedule(static,10)
  for (j=0; j<N; j++) {
    ... // some work here
  }
```



Keeps assigning chunks until done.


Chunk size that is not a multiple of the loop will result in thread with uneven numbers.

OpenMP: Problem with Static Scheduling



What happens if loop iterations do not take the same amount of time ?

Load imbalance



OpenMP: Dynamic Scheduling



Chunks are assigned on the fly, as threads become available
When a thread finishes on chunk, it is assigned another

```
#pragma omp parallel for default(shared) private(j) schedule(dynamic,10)
  for (j=0; j<N; j++) {
    ... // some work here
  }
```

Caveat Emptor: higher overhead than static!

OpenMP omp for Scheduling Recap



```
#pragma omp parallel for schedule(type [,size])
```

Scheduling types

Static

- Chunks of specified size assigned round-robin

Dynamic

- Chunks of specified size are assigned when thread finishes previous chunk

Guided

- Like dynamic, but chunks are exponentially decreasing
- Chunk will not be smaller than specified size

Runtime

- Type and chunk determined at runtime via environment variables

OpenMP: API



- API for library calls that perform useful functions
- Must include "omp:h"
- Will not compile without OpenMP compiler support

```
#include <omp.h> //<-- necessary header file for OpenMP API
#include <stdio.h>

int main(int argc, char *argv[]){

    printf("OpenMP running with %d threads\n", omp_get_max_threads());

    #pragma omp parallel
    {
        //Code here will be executed by all threads
        printf("Hello World from thread %d\n", omp_get_thread_num());
    }

    return 0;
}
```

OpenMP: API



```
void omp_set_num_threads(int num_threads)
```

Sets number of threads used in next parallel section
Overrides OMP_NUM_THREADS environment variable
Positive integer

```
int omp_get_max_threads()
```

Returns max possible (generally set by OMP_NUM_THREADS)

```
int omp_get_num_threads()
```

Returns number of threads currently in team

```
int omp_get_thread_num()
```

Returns thread id of calling thread
Between 0 and omp_get_num_threads-1


```
double omp_get_wtime()
```

Returns number of seconds since some point
Use in pairs time=(t2-t1)

OpenMP: Performance Tips



- Avoid serialization!
- Avoid using `#pragma omp parallel` for before each loop
 - Can have significant overhead
 - Thread creation and scheduling is NOT free!!
 - Try for broader parallelism
 - One `#pragma omp parallel`, multiple `#pragma omp for`
 - Always try to parallelize the outer most loop
- Use reduction whenever possible
- Minimize I/O
- Minimize `critical`
 - Use `atomic` instead of `critical` where possible





Hybrid OpenMP-MPI


Hybrid MPI and OpenMP

Hybrid application programs using **MPI + OpenMP** are now common place on large HPC systems.

There are basically two main motivations:

1. Reduced memory footprint, both in the application and in the MPI library (eg communication buffers).
2. Improved performance, especially at high core counts where pure MPI scalability runs out.

A common hybrid approach



- From dequential code, alongside MPI first, then try adding OpenMP
- From MPI code, add OpenMP
- From OpenMP code, treat as serial code
- The simplest and least error-prone method is to use MPI outside the parallel region and allow only the master thread to communicate between MPI tasks.
- Could use MPI in parallel region with thread-safe MPI.

Hybrid MPI and OpenMP

- Two-level Parallelization
 - Mimics hardware layout of cluster
 - Only place this really make sense
 - MPI between nodes
 - OpenMP within shared-memory nodes
- Why ?
 - Saves memory by not duplicating data
 - Minimize interconnect communication by only having 1 MPI process per node
- Careful of MPI calls within OpenMP block
 - Safest to do MPI calls outside (but not required)

Obviously requires some thought!

Hybrid Programming

In hybrid programming each process can have multiple threads executing simultaneously
All threads within a process share all MPI objects Communicators, requests, etc.

MPI defines 4 levels of thread safety:

- **MPI_THREAD_SINGLE** : One thread exists in program
- **MPI_THREAD_FUNNELED** : Multithreaded but only the master thread can make MPI calls Master is one that calls `MPI_Init_thread()`
- **MPI_THREAD_SERIALIZED**: Multithreaded, but only one thread can make MPI calls at a time
- **MPI_THREAD_MULTIPLE**: Multithreaded and any thread can make MPI calls at any time. Use `MPI_Init_thread` instead of `MPI_Init` if more than single thread

Hybrid Programming

Safest (easiest) to use `MPI_THREAD_FUNNLED`

- Fits nicely with most OpenMP models
 - - Expensive loops parallelized with OpenMP
 - - Communication and MPI calls between loops
- Eliminates need for true “thread-safe” MPI
- Parallel scaling efficiency may be limited (Amdahl’s law) by `MPI_THREAD_FUNNLED` approach
- Moving to `MPI_THREAD_MULTIPLE` does come at a performance price (and programming challenge)

Hybrid Programming Example

Program hybrid

```
call MPI_INIT (ierr)
call MPI_COMM_RANK (...)
call MPI_COMM_SIZE (...)
... some computation and MPI
communication
... start OpenMP within node
!$OMP PARALLEL DO PRIVATE(i)
!$OMP&                SHARED(n)
  do i=1,n
    ... computation
  enddo
!$OMP END PARALLEL DO
... some computation and MPI
communication
call MPI_FINALIZE (ierr)
end
```

Start with MPI Initialization

Create OMP parallel regions with MPI task (Process):

- Serial Regions are the master head or MPI task .
- MPI rank is known to all threads.

Call MPI library in serial and parallel regions.

Finalize MPI

Hybrid Programming Example

- Numerical integration

$$\int_0^1 \frac{4}{1+x^2} dx = \pi$$


- Discretization:

$$\Delta = 1/N: \text{step} = 1/\text{NBIN}$$

$$x_i = (i+0.5)\Delta \quad (i = 0, \dots, N-1)$$

$$\sum_{i=0}^{N-1} \frac{4}{1+x_i^2} \Delta \cong \pi$$

```
#include <stdio.h>
#define NBIN 100000
void main() {
    int i; double step,x,sum=0.0,pi;
    step = 1.0/NBIN;
    for (i=0; i<NBIN; i++) {
        x = (i+0.5)*step;
        sum += 4.0/(1.0+x*x);}
    pi = sum*step;
    printf("PI = %f\n",pi);
}
```



```
#include <stdio.h>
#include <stdlib.h>
#include <mpi.h> /* MPI header file */
#include <omp.h> /* OpenMP header file */
#define NUM_STEPS 100000000
#define MAX_THREADS 4

int main(int argc, char *argv[]) {
    int nprocs, myid;
    int tid, nthreads, nbin;
    double start_time, end_time;
    double pi, Psum=0.0, sum[MAX_THREADS]={0.0};
    double step = 1.0/(double) NUM_STEPS;

    /* initialize for MPI */
    MPI_Init(&argc, &argv); /* starts MPI */
    /* get number of processes */
    MPI_Comm_size(MPI_COMM_WORLD, &nprocs);
    /* get this process's number (ranges from 0 to nprocs - 1) */
    MPI_Comm_rank(MPI_COMM_WORLD, &myid);

    nbin= NUM_STEPS/nprocs;
```


Hybrid Programming Example

```
#pragma omp parallel private(tid)
{
    int i;
    double x;
    nthreads=omp_get_num_threads();
    tid=omp_get_thread_num();
    for (i=nbin*myid+tid; i < nbin*(myid+1); i+= nthreads) { /* changed */
        x = (i+0.5)*step;
        sum[tid] += 4.0/(1.0+x*x);
    }
}
for(tid=0; tid<nthreads; tid++) /*sum by each mpi process*/
    Psum += sum[tid]*step;

MPI_Reduce(&Psum, &pi, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);/* added */

if (myid == 0) {
    printf("parallel program results with %d processes:\n", nprocs);
    printf("pi = %g (%17.15f)\n", pi, pi);
}
MPI_Finalize();

return 0;
}
```

Results

- **MPI**
MPI uses 8 processes:
pi = 3.14159 (3.141592653589828)
- **OpenMP**
OpenMP uses 8 threads:
pi = 3.14159 (3.141592653589882)
- **Hybrid**
mpi process 0 uses 4 threads
mpi process 1 uses 4 threads
mpi process 1 sum is 1.287 (1.287002217586605)
mpi process 0 sum is 1.85459 (1.854590436003132)
Total MPI processes are 2
pi = 3.14159 (3.141592653589738)



Thank you for your attention !

