



# **P**ARALLEL **P**ROGRAMMING...

Copyright 2023 Patrick Lemoine. All rights reserved.

# Parallel Programming: Overview

SESSION 3/6



**Programming Interface for parallel computing**

OpenMP (Open Multi-Processing)

# **P**rogramming **i**nterface for **p**arallel **c**omputing

병렬 컴퓨팅을 위한 프로그래밍  
인터페이스



OpenMP (**O**pen **M**ulti-**P**rocessing)

OpenMP®




# OpenMP (Open Multi-Processing)

**OpenMP** is a programming interface for parallel computing on **shared memory architecture**.

A set of compiler directives and a runtime library

```
#pragma omp parallel num_threads(4)
```

```
#include <omp.h>
```



- It allows you to manage:





- the creation of light processes
- the sharing of work between these lightweight processes
- synchronizations (explicit or implicit) between all light processes
- the status of the variables (private or shared).

# OpenMP (Open Multi-Processing)

## OpenMP is based on Fork/Join model

- When program starts, one Master thread is created
- Master thread executes sequential portions of the program
- At the beginning of parallel region, master thread forks new threads
- All the threads together now forms a “team”
- **At the end of the parallel region, the forked threads die !**



# OpenMP (Open Multi-Processing)

## The OpenMP API consists of

- compiler directives (for insertion *into sequential* Fortran/C/C++\$code)
- a few library routines
- some environment variables



### Advantages:

- User-friendly
- Incremental parallelization of a serial code
- Possible to have a single source code for both serial and parallelized versions




### Disadvantages:

- Relatively limited user control
- Most suitable for parallelizing loops (data parallelism)
- Performance?

# OpenMP (Open Multi-Processing)

## What is a **Shared-Memory Program**?

- One process that spawns multiple threads
- Threads can communicate via shared memory
  - Read/Write to shared variables
  - Synchronization can be required!
- OS decides how to schedule threads





# OpenMP (Open Multi-Processing)

- **Shared memory model**

- Threads communicate by accessing shared variables.

- **The sharing is defined syntactically**

- Any variable that is seen by two or more threads is shared.
- Any variable that is seen by one thread only is private.



- **Race conditions possible**

- Use synchronization to protect from conflicts.
- Change how data is stored to minimize the synchronization.

# OpenMP (Open Multi-Processing): Why ?

- **Multicore CPUs are everywhere:**

- Servers with over 100 cores today and more.
- Even smartphone CPUs have 8 cores.

- **Multithreading, natural programming model**

- All processors share the same memory.
- Threads in a process see same address space.
- Many shared-memory algorithms developed.



- **Multithreading is hard**

- Lots of expertise necessary.
- Deadlocks and race conditions.
- **Non-deterministic** behavior makes it hard to debug.


# OpenMP: Process and thread: what is the difference ?



- You need an **existing process** to **create a thread**.
- **Each process** has at least **one thread of execution**.
- A **process** has its **own virtual memory space** that **cannot be accessed by other processes running** on the same or on a different processor.
- **All threads created** by a process **share the virtual address space of that process**.
  - They read and write to the same address space in memory.
  - They share the same process and user ids, file descriptors, and signal handlers.
  - They have their own program counter value and stack pointer, and can run independently on several processors.

# OpenMP: Terminology and behavior

- **OpenMP Team** = **Master** + **Worker**
- **Parallel Region** is a block of code executed by all threads simultaneously (*has implicit barrier*)
  - The master thread always has thread id **0** !
  - Parallel regions can be nested.
  - If **clause** can be used to **guard the parallel region**.



# OpenMP: General Code Structure



```
#include <omp.h>
```

```
main ()  
{
```


```
    int var1, var2, var3;  
    //serial code
```

```
    //start of a parallel region
```

```
    #pragma omp parallel private(var1, var2) shared(var3)  
    {...}
```

```
    //more serial code  
    {...}
```

```
    //another parallel region  
    #pragma omp parallel  
    {...}  
}
```



# OpenMP: Constructs Parallel Region



## Parallel region

- Thread creates team, and becomes master (id 0).
- All threads run code after.
- Barrier at end of parallel section.



```
#pragma omp parallel [clause ...]  
    if (scalar_expression)  
    private (list)  
    shared (list)  
    default (shared | none)  
    firstprivate (list)  
    lastprivate (list)  
    reduction (operator: list)  
    num_threads (integer)
```

*structured\_block*

(not a complete list)

# OpenMP: Parallel Clauses



```
#pragma omp parallel if (scalar_expression)
```

Only execute in parallel.  
Otherwise serial.

```
#pragma omp parallel private (list)
```

Data local to thread.  
Value are **not guaranteed to be defined on exit** (even if defined before)  
No storage associated with original object  
*Use firstprivate and/or lastprivate clause to override*

```
#pragma omp parallel firstprivate (list)
```

Variables in list are private.  
Initialized with the value the variable had before entering the construct.

```
#pragma omp parallel for lastprivate (list)
```

Only in for loops  
Variables in list are private.  
The thread that executes the sequentially last iteration updates the value of the variables in the list.



# OpenMP: Parallel Clauses



```
#pragma omp shared (list)
```

Data is accessible by all threads in team.  
All threads access same address space.

Improperly scoped variables are big source of OMP bugs

- Shared when should be private
- Race condition

```
#pragma omp default (shared | none)
```


Tip: *Safest is to use default (none) and declare by hand.*



# OpenMP: Barrier



- When a thread reaches a barrier it only continues after all the threads in the same thread team have reached it.
- Each barrier must be encountered by all threads in a team, or none at all
- The sequence of work-sharing regions and barrier regions encountered must be same for all threads in team
- Implicit barrier at the end of: `do`, `parallel`, `single`, `workshare`




# OpenMP: Caution Race Condition



When multiple threads simultaneously read/write  
Multiple OMP solutions :

- Reduction
- Atomic
- Critical

```
#pragma omp parallel for private(i) shared(sum)
for (i=0; i<N; i++) {
    sum += i;
}
```



Should be 3!

# OpenMP: Critical Section



One solution: use **critical**


Only one thread at a time can execute a critical section

```
#pragma omp critical
{
    sum += i;
}
```

Downside ?

YES SLOOOOWWW

Overhead and serialization




# OpenMP: Atomic



**Atomics** like "mini" critical  
Only one line  
Certain limitations

```
#pragma omp atomic  
sum += i;
```

*Hardware controlled*  
*Less overhead the critical*



# OpenMP: Reduction



```
#pragma omp reduction (operator:variable)
```

- **Avoids race condition**
- Reduce variable must be shared
- Makes variable private, then performs operator at end of loop
- operator cannot be overloaded (c++)

One of: +, \*, -, / (and &, ^, |, &&, ||)

OpenMP 3.1: added min and max for c/c++

```
#include <omp.h>
#include <stdio.h>

int main() {

    int i;
    const int N = 1000;
    int sum = 0;

    #pragma omp parallel for private(i) reduction(+: sum)
    for (i=0; i<N; i++) {
        sum += i;
    }

    printf("reduction sum=%d (expected %d)\n", sum, ((N-1)*N)/2);
}
```

# OpenMP: Scheduling omp for



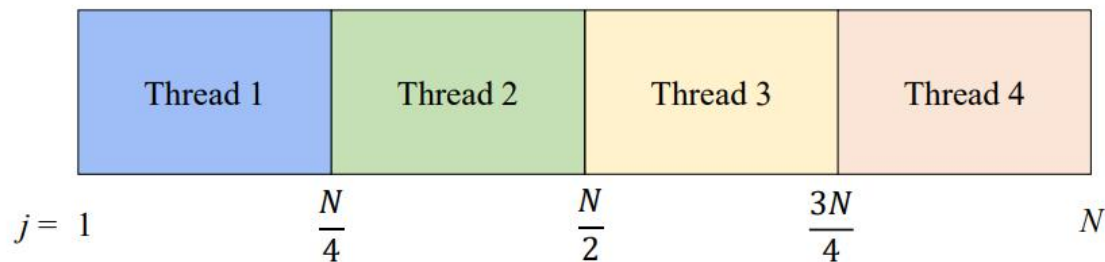
How does a loop get split up ? *In **MPI**, we have to do it manually!!!*

**If you do not tell what to do, the compiler decides**

**Usually compiler chooses "static" - chunks of  $N/p$**

```
#pragma omp parallel for default(shared) private(j)
  for (j=0; j<N; j++) {
    ... // some work here
  }
```

Unspecified schedule




# OpenMP: Static Scheduling



You can tell the compiler what size chunks to take ?

```
#pragma omp parallel for default(shared) private(j) schedule(static,10)
for (j=0; j<N; j++) {
    ... // some work here
}
```



Keeps assigning chunks until done.

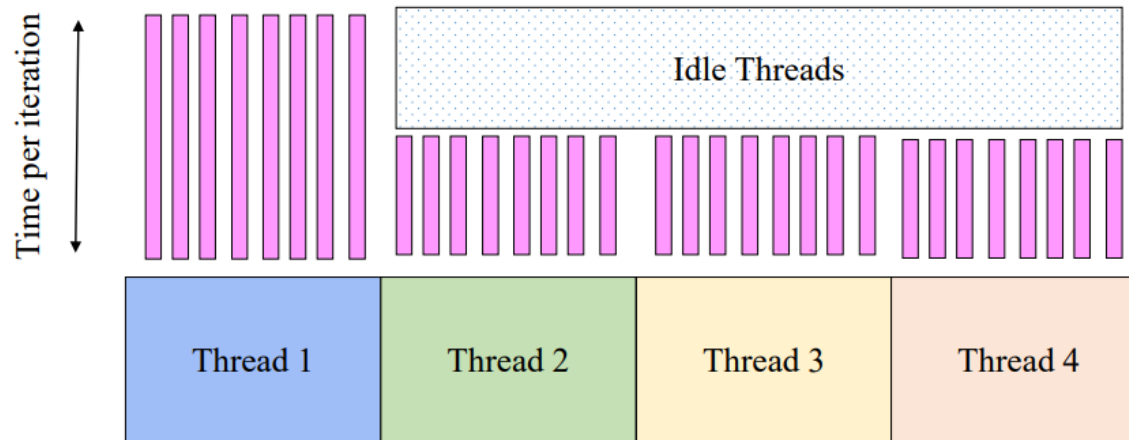
**Chunk size that is not a multiple of the loop** will result in thread with uneven numbers.

# OpenMP: Problem with Static Scheduling



What happens if loop iterations do not take the same amount of time ?

**Load imbalance**





# OpenMP: Dynamic Scheduling



Chunks are assigned on the fly, as threads become available.

When a thread finishes on chunk, it is assigned another

```
#pragma omp parallel for default(shared) private(j) schedule(dynamic,10)
for (j=0; j<N; j++) {
    ... // some work here
}
```

*Caveat: higher overhead than static!*

# OpenMP: For Scheduling Recap



```
#pragma omp parallel for schedule(type [,size])
```

## Scheduling types

### Static

- Chunks of specified size assigned round-robin

### Dynamic

- Chunks of specified size are assigned when thread finishes previous chunk

### Guided

- Like dynamic, but chunks are exponentially decreasing
- Chunk will not be smaller than specified size

### Runtime

- Type and chunk determined at runtime via environment variables

# OpenMP: API



- API for library calls that perform useful functions
- Must include "omp:h"
- Will not compile without OpenMP compiler support

```
#include <omp.h> //<-- necessary header file for OpenMP API
#include <stdio.h>

int main(int argc, char *argv[]){

    printf("OpenMP running with %d threads\n", omp_get_max_threads());

    #pragma omp parallel
    {
        //Code here will be executed by all threads
        printf("Hello World from thread %d\n", omp_get_thread_num());
    }

    return 0;
}
```

# OpenMP: API



```
void omp_set_num_threads(int num_threads)
```

Sets number of threads used in next parallel section  
Overrides OMP\_NUM\_THREADS environment variable

```
int omp_get_max_threads()
```

Returns max possible (generally set by OMP\_NUM\_THREADS)

```
int omp_get_num_threads()
```

Returns number of threads currently in team

```
int omp_get_thread_num()
```

Returns thread id of calling thread  
Between 0 and omp\_get\_num\_threads-1

```
double omp_get_wtime()
```

Returns number of seconds since some point  
Use in pairs time=(t2-t1)

# OpenMP: Example Fibonacci



```
int fib ( int n )
{
    int x,y;
    if ( n < 2 ) return n;
    #pragma omp task shared(x,n)
    x = fib(n-1);


    #pragma omp task shared(y,n)
    y = fib(n-2);

    #pragma omp taskwait
    return x+y;
}
```

```
int main()
{
    #pragma omp parallel
    #pragma omp single nowait
    result = comp_fib_numbers(10);
    return EXIT_SUCCESS;
}
```

$\text{fib}(0) = 1$   
 $\text{fib}(1) = 1$   
 $\text{fib}(n) = \text{fib}(n-1) + \text{fib}(n-2)$

avec  $n \in \mathbb{N}$



# OpenMP: Example Quicksort

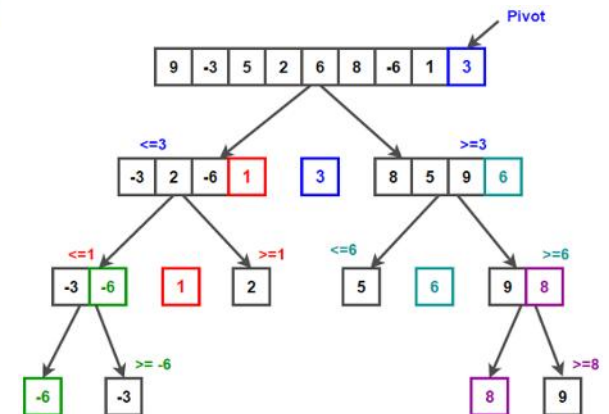


```
void quick_sort (int p, int r, float *data)
{
    if (p < r) {
        int q = partition (p, r, data);
        #pragma omp task

        quick_sort (p, q-1, data, low_limit);
        #pragma omp taskwait

        quick_sort (q+1, r, data, low_limit);
    }
}
```

```
void par_quick_sort (int n, float *data)
{
    #pragma omp parallel
    {
        #pragma omp single nowait
        quick_sort (0, n, data);
    }
}
```




# OpenMP: Example Cholesky Factorization

The **Cholesky factorization**, also known as Cholesky decomposition, is a process of breaking down of a Hermitian, positive-definite matrix into the product of a lower triangular matrix and its conjugate transpose, which is important for quick numerical solutions in linear algebra.

- 1: Input: Dictionary  $\mathbf{D}$ , signal  $\underline{x}$ , target sparsity  $K$  or target error  $\epsilon$
- 2: Output: Sparse representation  $\underline{\gamma}$  such that  $\underline{x} \approx \mathbf{D}\underline{\gamma}$
- 3: Init: Set  $I := ()$ ,  $\mathbf{L} := [1]$ ,  $\underline{r} := \underline{x}$ ,  $\underline{\gamma} := \underline{0}$ ,  $\underline{\alpha} := \mathbf{D}^T \underline{x}$ ,  $n := 1$
- 4: **while** (*stopping criterion not met*) **do**
- 5:    $\hat{k} := \underset{k}{\text{Argmax}} |d_k^T \underline{r}|$
- 6:   **if**  $n > 1$  **then**
- 7:      $\underline{w} := \text{Solve for } \underline{w} \{ \mathbf{L}\underline{w} = \mathbf{D}_I^T d_{\hat{k}} \}$
- 8:      $\mathbf{L} := \begin{bmatrix} \mathbf{L} & \underline{0} \\ \underline{w}^T & \sqrt{1 - \underline{w}^T \underline{w}} \end{bmatrix}$
- 9:   **end if**
- 10:    $I := (I, \hat{k})$
- 11:    $\underline{\gamma}_I := \text{Solve for } \underline{c} \{ \mathbf{L}\underline{L}^T \underline{c} = \underline{\alpha}_I \}$
- 12:    $\underline{r} := \underline{x} - \mathbf{D}_I \underline{\gamma}_I$
- 13:    $n := n + 1$
- 14: **end while**

```
void cholesky(int ts, int nt, double* a[nt][nt]) {  
    for (int k = 0; k < nt; k++) {  
        // Diagonal Block factorization  
        #pragma omp task depend(inout: a[k][k])  
        potrf(a[k][k], ts, ts);  
  
        // Triangular systems  
        for (int i = k + 1; i < nt; i++) {  
            #pragma omp task depend(in: a[k][k])  
                        depend(inout: a[k][i])  
            trsm(a[k][k], a[k][i], ts, ts);  
        }  
  
        // Update trailing matrix  
        for (int i = k + 1; i < nt; i++) {  
            for (int j = k + 1; j < i; j++) {  
                #pragma omp task depend(inout: a[j][i])  
                            depend(in: a[k][i], a[k][j])  
                dgemm(a[k][i], a[k][j], a[j][i], ts, ts);  
            }  
            #pragma omp task depend(inout: a[i][i])  
                        depend(in: a[k][i])  
            syrk(a[k][i], a[i][i], ts, ts);  
        }  
    }  
}
```




# OpenMP: Example Gauss-Seidel

```
void gauss_seidel(int tsteps, int size, int TS, int (*p)[size]) {
    int NB = size / TS;

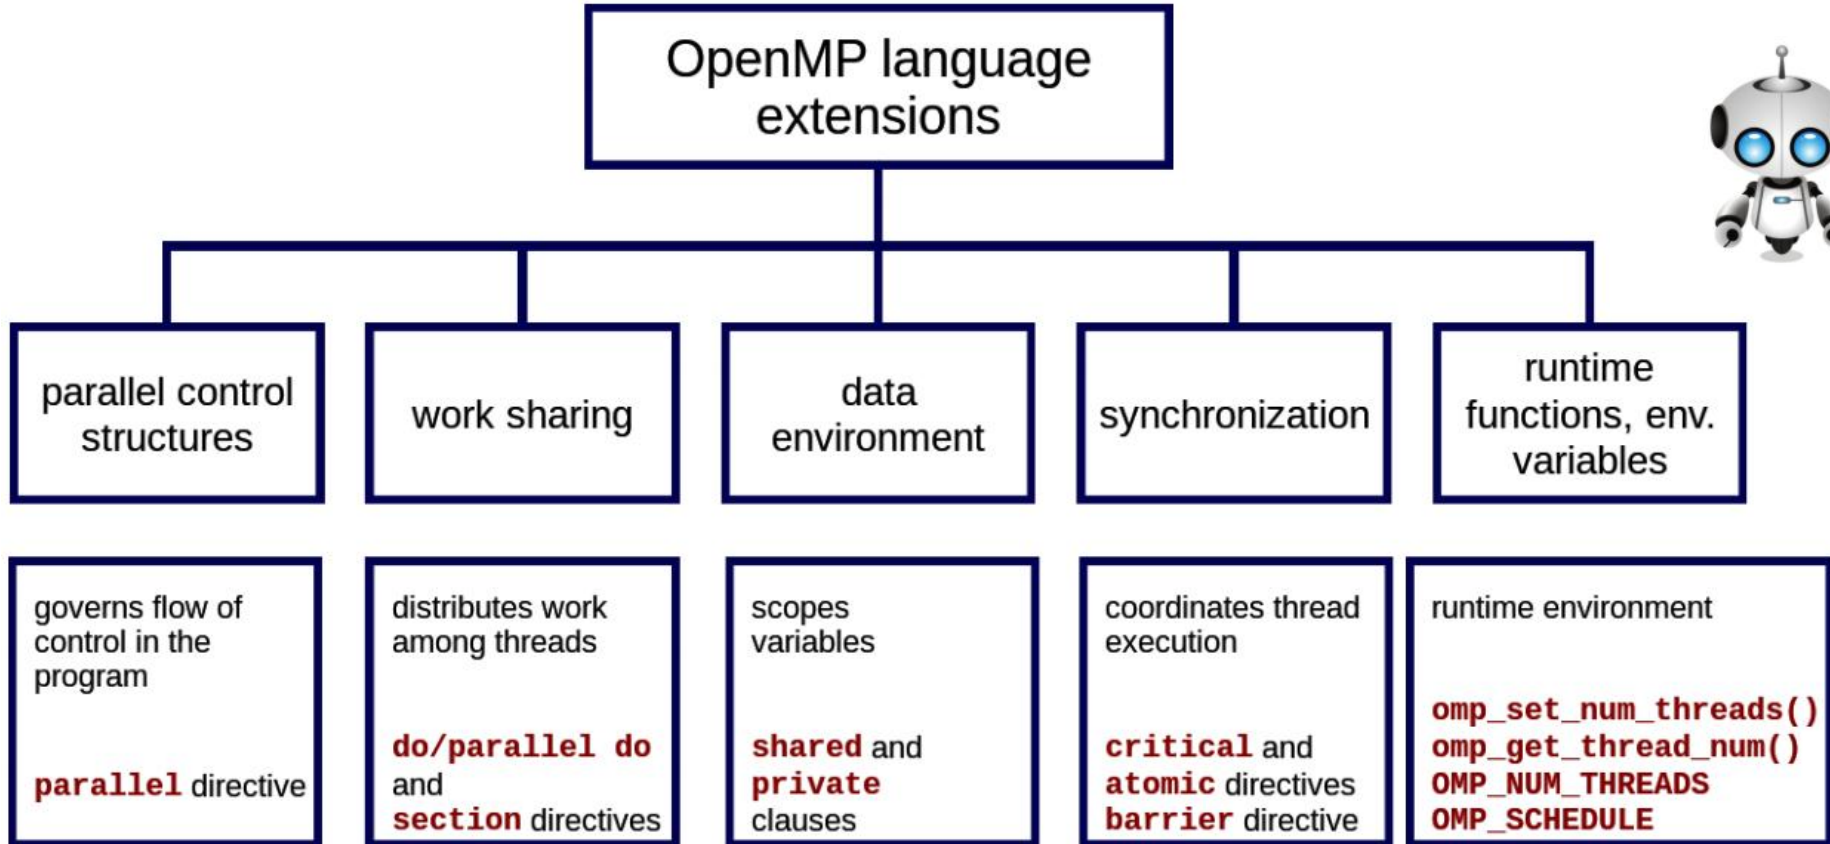
    #pragma omp parallel
    #pragma omp single
    for (int t = 0; t < tsteps; ++t)
        for (int ii=1; ii < size-1; ii+=TS)
            for (int jj=1; jj < size-1; jj+=TS) {
                #pragma omp task depend(inout: p[ii:TS][jj:TS])
                depend(in: p[ii-TS:TS][jj:TS], p[ii+TS:TS][jj:TS],
                      p[ii:TS][jj-TS:TS], p[ii:TS][jj+TS:TS])
                {
                    for (int i=ii; i<(1+ii)*TS; ++i)
                        for (int j=jj; j<(1+jj)*TS; ++j)
                            p[i][j] = 0.25 * (p[i][j-1] * p[i][j+1] *
                                                  p[i-1][j] * p[i+1][j]);
                }
            }
}
```

**Gauss-Seidel Method** is used to solve the linear system Equations. It is a method of iteration for solving  $n$  linear equation  $Ax=b$  with the unknown variables.





# OpenMP: Summary Overview



# OpenMP: Performance Tips...



- Avoid serialization !
- Avoid using **#pragma omp parallel** for before loop
  - Can have significant overhead
    - Thread creation and scheduling is not free!!
  - Try for broader parallelism
    - One **#pragma omp parallel**, multiple **#pragma omp for**
  - Always try to parallelize the other most loop
- Use reduction whenever possible
- Minimize I/O
- Minimize critical
  - Use **atomic** instead of critical where possible



Thank you for your attention !

